FOR CARRIERS AND THE AI VENDORS WHO SERVE THEM

The minimum AI decision record

Ten questions for an insurance AI vendor, per decision

Pick one real decision the AI touched. Answer from the record itself, not from what the vendor could assemble for an exam. Every no is a gap.

1	For one specific decision, can you show what data the AI saw?	Inputs
2	Can you show what it recommended or decided, and why?	Decision and reasons
3	Which policy provisions or rules did it rely on, and which version was in force that day?	Policy and version
4	Which model and version answered, and would the record show if it changed later?	Model and version
5	Who reviewed it, and did they agree, change or override the AI?	Human review
6	When did each step happen?	Time
7	Was the record made at the time of the decision, not assembled later for an exam?	Contemporaneous
8	Can the carrier open it without logging into the vendor's system?	Portable
9	Would the carrier notice if the record were changed after it was handed over?	Tamper-evident
10	Can you list decisions that should have a record but don't?	Coverage

How to score: one point per question you can answer from the record itself. 9–10 a decision record a carrier can rely on · 6–8 close, but reasons, review or timing are incomplete · 3–5 observability data · 0–2 logs. **Every no is a gap.**

What it does not ask for. None of these ten questions needs the vendor's source code, training data or model weights. They ask what happened in one decision and whether a person can read it.

How to use it. Take one real decision the AI touched. If you cannot answer a question from the record, that is the gap to close before an examiner asks. A worked example, with invented claim data and a real override, is at epilabs.org/insurance/ai-vendor-decision-record.